The σ_T Variance Proxy and Its ε Inversion: A Unified Mechanism Across Deterministic and Stochastic FX Dynamics

with a cadCAD-wrapping Bridge to the AI-Cost-Factor Downstream Consumer

Juan Manuel Serrano
juan.serranotmf@gmail.com

2026-05-16 (working paper draft)

Abstract

We close notes/PRIMITIVES.md §15 open item 2 (*stochastic-FX variant — Path A v3*) by verifying, across three SDE families (geometric Brownian motion, Ornstein–Uhlenbeck, and Merton jump-diffusion), that the realized-variance proxy $\sigma_T = T^{-1} \sum_j (X_j - \bar{X}_{\text{path}})^2$ admits a single family-agnostic inversion $\varepsilon(\sigma_T) = \sqrt{8\,\sigma_T/\bar{x}^2}$ that holds pointwise (PIN Z1.3A, algebraic identity) and distributionally (PIN Z1.3B mean-only moment match against literature-derived discrete-statistic analytic moments; PIN Z1.4 Kolmogorov–Smirnov goodness-of-fit against per-family-dispatched reference distributions). Phase B (moment) and Phase C (KS) verdicts pass at canonical pin under each SDE family at a fixed $N = 1000$ Monte-Carlo budget (PIN Z1.5 anti-fishing floor); audit-block hash determinism (PIN Z1.2) is bit-exact across both intra-process and inter-process re-runs.

The unified mechanism is wrapped, in a forthcoming integration, into a cadCAD dynamic-agent simulation where the verified σ_T machinery becomes the structural M-side substrate for downstream (Y, M, X)-triple iterations. We position the AI-cost-factor model (docs/specs/2026-05-16-ai-cost-factor-model-design.md v0.2.1, under construction) as the immediate downstream consumer: its candidate-X identification step is independent of the verification this paper provides, but its eventual M-design step inherits the σ_T-replicating Carr–Madan strip (cohort_5_strip, deterministic instantiation) and the per-family stochastic verification floor.

The paper makes no causal claim about the wage-to-productive-capital ratchet that motivates the broader Abrigo operating framework; the contribution is strictly the verification floor on the σ_T mechanism that any downstream (Y, M, X) instrument family must clear before its M-design step is admissible.

Contents

1 Unified framework: σ_T as a structural primitive

1.1 Operating context and the (Y, M, X)-triple unit of work

The Abrigo operating framework, summarized in `CLAUDE.md`, takes as its highest-level goal the minimization of income inequality framed in post-Keynesian terms (distribution as institutionally determined rather than equilibrium-given). The product family pursued by the project is *permissionless on-chain perpetual convex instruments* settled on Panoptic, with denomination chosen per target population (Mento-native COPm/BRLm/KESm/EURm/USDm, USDC, ETH, or sectoral baskets). The transmission channel from wage to productive capital runs through a *premium-funded ratchet*: a wage earner pays a small recurring premium out of wage income, and the instrument's accumulated convex payoff and roll yield convert over time into productive-capital exposure that absent the instrument the wage earner would never cross the wage–capital boundary to obtain.

The operating unit of work is the (Y, M, X)-triple:

- Y — outcome variable measuring a target population's exposure to the wage-to-capital transition;
- M — the Panoptic pool configuration hosting the hedge (underlying pair, strike/range geometry, payoff shape);
- X — the major risk currently blocking that transition for the target population.

This paper does not introduce a new (Y, M, X)-triple. It verifies a *shared substrate* that every (Y, M, X)-triple containing a convex on-chain hedge inherits.

1.2 The σ_T proxy and the ε inversion

Definition 1.1 (Realized-variance proxy, σ_T)**.** *Let $\{X_j\}_{j=0}^{n}$ denote a discrete FX-rate path observed at $n+1$ equally spaced timepoints $t_j = j \cdot \Delta t$ on the trading horizon $[0, T]$, with $\Delta t = T/n$ and $\bar{X}_{path} = (n+1)^{-1} \sum_{j=0}^{n} X_j$ the path's sample mean. The discrete realized-variance proxy is*

$$\sigma_T \;\coloneqq\; \frac{1}{T} \sum_{j=0}^{n} \big(X_j - \bar{X}_{path}\big)^2. \tag{1}$$

Equation (1) is the discrete-time form of notes/PRIMITIVES.md §6 eq. (7). The normalization is by T (not by $n+1$): the statistic carries units of $[X]^2$, and it scales proportionally to grid density. This is intentional and load-bearing: under any continuous-time FX-path generator, the sum-of-squared deviations sample more variation as $n \to \infty$, and dividing by T preserves comparability across grids at fixed T.

Definition 1.2 (ε inversion)**.** *Given a realized-variance proxy σ_T and an FX mean $\bar{x}$ (typically the canonical pin $\bar{x} = \bar{X}/\bar{Y} = 4000$ for the COP/USD-denominated cohort), the inversion is*

$$\varepsilon(\sigma_T) \;=\; \sqrt{\frac{8\,\sigma_T}{\bar{x}^2}}. \tag{2}$$

Equation (2) is PRIMITIVES.md §6 eq. (8). It is purely algebraic: no SDE-family parameter enters; the relation is family-agnostic. The square-root form ensures (PIN Z1.3A trivial-degenerate limit) that $\lim_{\sigma_T \to 0^+} \varepsilon(\sigma_T) \;=\; 0$, which captures the geometric fact that a flat (deterministic) path has zero implied volatility per unit FX mean.

1.3 Verification structure: three phases, six pins

The verification this paper reports has three phases, each anchored to a distinct numbered pin in docs/specs/2026-05-11-stochastic-fx-variant-design.md v0.5:

Phase A — algebraic identity (PIN Z1.3A). For each path's $\sigma_T^{(n)}$, the pointwise identity $\big(\varepsilon(\sigma_T^{(n)})\big)^2 = 8\sigma_T^{(n)}/\bar{x}^2$ holds to float64 precision ($\textsc{numerical\_identity\_tol} = 10^{-6}$). Family-agnostic; reduces to a re-evaluation of (2).

Phase B — mean-only moment match (PIN Z1.3B, spec v0.5 amendment). Empirical mean of $\sigma_T^{(n)}$ across the N-path ensemble matches the hand-derived analytic $\mathbb{E}[\sigma_T]$ for the DISCRETE statistic (1) at the canonical grid, within $\textsc{moment\_rel\_tol} = 0.05$. The discrete-statistic analytic uses the centering-projection identity $\mathbb{E}[\sigma_T \cdot T] = \operatorname{tr}\big(\mathcal{M}\,\Sigma\big) + \mu^\top \mathcal{M}\,\mu$ where $\mathcal{M} = I_{n+1} - (n+1)^{-1}\mathbf{1}\mathbf{1}^\top$ and Σ is the SDE-family-specific auto-covariance kernel on the canonical grid. Variance is NOT in this gate (see Remark 1.1).

Phase C — Kolmogorov–Smirnov goodness-of-fit (PIN Z1.4, spec v0.5 amendment). Per-family dispatch on the reference distribution shape (not on the threshold):

- GBM: method-of-moments lognormal reference fit to the analytic $(\mathbb{E}, \mathrm{Var})$, tested via `scipy.stats.ks_1samp`.
- OU: method-of-moments gamma reference fit to the analytic $(\mathbb{E}, \mathrm{Var})$, tested via `ks_1samp`.
- Merton: empirical-CDF reference via a high-N ($N_{\text{REF}} = 100\,000$) reference run from the same SDE, tested via `ks_2samp`. This avoids the shape mismatch between Merton's Poisson-mixture-of-lognormals geometry and the 2-parameter lognormal alternative.

Single threshold across all three families: $\textsc{ks\_pvalue\_floor} = 0.01$.

Remark 1.1 (Why mean-only Phase B; what variance discloses)**.** *At the PIN Z1.5 anti-fishing path-count floor $N = 1000$, the standard error of the sample-variance estimator on σ_T is intrinsically $\sqrt{(\kappa_{\mathit{eff}} - 1)/(N-1)}$ where κ_{eff} ranges 4–35 across the three families (driven by jump kurtosis for Merton). This produces an 8%–30% Monte-Carlo noise floor on the variance estimator — strictly incompatible with $\textsc{moment\_rel\_tol} = 0.05$. The spec v0.4 attempt to gate variance against the same 5% tolerance was mathematically unsatisfiable; v0.5 §11.6 dropped variance from the Phase B gate. Variance relative error is still computed and emitted on the InversionVerdict as an audit-trail observation (it is not zero); the full-distribution match is preserved at Phase C, where the KS test against a moment-matched reference constrains the joint shape (mean, variance, skewness, kurtosis) via the empirical CDF.*

1.4 Pin coverage and anti-fishing posture

The six pins anchoring this paper are:

Pin	Statement
PIN Z1.1	SDE parameters pre-pinned ex-ante per family (§3, §4, §5 below).
PIN Z1.2	Path-ensemble determinism: identical $(\text{params}, \text{rng\_seed}, n_{\text{paths}})$ triples yield bit-exact paths and bit-exact audit_block.
PIN Z1.3A	Algebraic identity $\big(\varepsilon(\sigma_T)\big)^2 = 8\sigma_T/\bar{x}^2$ to float64 precision.
PIN Z1.3B	Mean-only moment match against discrete-statistic analytic $\mathbb{E}[\sigma_T]$ within 5%.
PIN Z1.4	Per-family Phase C KS p-value ≥ 0.01.
PIN Z1.5	Anti-fishing routing: failures route to CORRECTIONS-α and scoped two-reviewer re-review; NEVER silent parameter re-tuning; NEVER "increase N until passing"; $N = 1000$ is an invariant, not a parameter to grow.
PIN Z1.6	Strip preservation: the cohort_5_strip/IronCondor_strip.json audit_block remains byte-equal across the verification plan's execution window.

PIN Z1.5 carries forward the broader project's anti-fishing invariants ($N_{\text{min}} = 75$, $\text{POWER}_{\text{min}} = 0.80$, MDES_SD = 0.40 SD-units of Y, pre-pinned sign expectation and lag structure). For this paper specifically, three concrete anti-fishing prohibitions hold across all four model instantiations: (i) no silent threshold relaxation, (ii) no seed re-roll on failure, (iii) no N-floor lowering. The four spec amendments (§11.6, §11.7 of the parent spec; §16.3, §16.4, §16.5, §16.7 of the plan) are each subjected to scoped two-reviewer re-review per master-spec §6.4, and each preserves these three prohibitions.

1.5 Downstream consumer: the AI-cost-factor model

The verification floor established by this paper is the immediate prerequisite for an (Y, M, X)-triple targeting the COP-denominated AI-tooling cost burden borne by non-subscribed LATAM developers (the AI-cost-factor model under construction, spec v0.2.1; see §6.3). That model's X-channel identification step (R5 descriptive risk quantification + R4-Stage-3 vol-clustering regression) operates on the empirical cost path, NOT on σ_T directly. But should X-identification yield a candidate FX-vol risk that admits Panoptic-eligible M-design, the M-side instrument is a σ_T-replicating strip and inherits the per-family stochastic verification floor of this paper.

This positioning is asymmetric on purpose: the σ_T machinery is *substrate* for downstream (Y, M, X)-triple iterations, not the final product. The contribution of this paper is the substrate verification; downstream consumers carry their own identification burdens.

2 Deterministic instantiation: Carr–Madan replication

2.1 Deterministic FX-path generator and its σ_T limit

The deterministic instantiation of the framework specializes the FX trajectory $\{X_j\}_{j=0}^{n}$ of Definition 1.1 to a closed-form perturbation around the cohort's stationary mean $\bar{X}/\bar{Y}$. Following notes/PRIMITIVES.md §5 eq. (6), the generator is

$$(X/Y)_t(\varepsilon,\omega) \;=\; \Big(1+\varepsilon\big(\cos^2(\omega t)-\tfrac{1}{2}\big)\Big)\,\frac{\bar{X}}{\bar{Y}}, \qquad 0<\varepsilon<1, \tag{3}$$

with perturbation kernel $f_t \coloneqq \cos^2(\omega t)-\frac{1}{2}$ so that $(X/Y)_t = (1+\varepsilon f_t)\,\bar{X}/\bar{Y}$. The amplitude ε controls the path's oscillation magnitude around the stationary mean and the frequency ω controls its rate. The deterministic generator is the noiseless limit of the stochastic SDE families instantiated in §3–5: it admits a closed-form σ_T and so anchors the algebraic identity $\big(\varepsilon(\sigma_T)\big)^2 = 8\sigma_T/\bar{x}^2$ of equation (2) at $\bar{X} = \bar{X}$, $\bar{Y} = \bar{Y}$, $\bar{x} = \bar{X}/\bar{Y}$.

2.2 The σ_0 closed-form anchor

Proposition 2.1 (Deterministic-limit reduction)**.** *Under generator (3) the realized-variance proxy σ_T of Definition 1.1 reduces, in the long-horizon ergodic average over the kernel f_t, to the closed form*

$$\sigma_0 \;=\; \frac{(\bar{X}/\bar{Y})^2\,\varepsilon^2}{8}, \tag{4}$$

yielding the inverse map $\varepsilon = \sqrt{8\sigma_0/(\bar{X}/\bar{Y})^2}$ that is the deterministic specialization of the family-agnostic inversion (2).

Equation (4) is the R1 anchor of notes/STAGE_2_RESULTS.md §2.1; it is obtained by substituting $f_t = \cos^2(\omega t) - \frac{1}{2}$ into the squared-deviation sum of equation (1) and exploiting $\langle \cos^2(\omega t) - \frac{1}{2}\rangle_t = 0$, $\langle\big(\cos^2(\omega t) - \frac{1}{2}\big)^2\rangle_t = 1/8$ on a horizon that resolves the period $2\pi/\omega$. At the canonical cohort pin $\bar{X}/\bar{Y} = 4000$, $\varepsilon = 0.10$, identity (4) pins $\sigma_0 = 2{,}000$; equation (2) recovers $\varepsilon(\sigma_T) = 0.10$ exactly.

2.3 Carr–Madan static replication

The Carr–Madan static-replication theorem [4] establishes that any twice-differentiable European payoff $g(S_T)$ admits an exact representation as a portfolio of vanilla out-of-the-money puts and calls, plus a bond and a forward leg:

$$g(S_T) \;=\; g(S_0) + g'(S_0)(S_T - S_0) + \int_0^{S_0} g''(K)\,(K - S_T)^+\,dK + \int_{S_0}^{\infty} g''(K)\,(S_T - K)^+\,dK. \tag{5}$$

The canonical Carr–Madan variance contract is the log-contract, $g(S_T) = -2\log(S_T/S_0) + 2(S_T - S_0)/S_0$, for which $g''(K) = 2/K^2 > 0$ — the canonical convex weight profile that, when integrated against the option-strike continuum (5), returns σ_T at expiry up to the usual martingale correction.

Remark 2.1 (From integral to discrete strip)**.** *The continuous representation (5) is not Panoptic-eligible as written; Panoptic's pool mechanics admit only finite-leg long-vol positions on Uniswap v3/v4 ranges. The deployment-eligible discretization replaces the two strike integrals in (5) by a finite sum of IronCondor cells, each cell contributing a piecewise-linear convex bump on a strike window. The 12-leg three-condor strip introduced in §2.4 is the smallest Panoptic-feasible discretization that preserves long-vol convexity throughout the convex span.*

2.4 The 12-leg IronCondor strip geometry

The replicating instrument is a three-condor strip stacked at progressively wider strikes around the cohort's stationary spot $S_0 = \bar{X}/\bar{Y}$. Each condor contributes four legs in the long-vol convention (*reverse* IronCondor: short the outermost strikes, long the inner strikes):

$$P_{\text{cond}}(S_T; K_1, K_2, K_3, K_4) \;=\; -(K_1 - S_T)^+ + (K_2 - S_T)^+ + (S_T - K_3)^+ - (S_T - K_4)^+, \tag{6}$$

with strikes $K_1 < K_2 < K_3 < K_4$. The strip payoff is the weighted sum

$$\Pi_{\text{strip}}(S_T) \;=\; \sum_{j=1}^{3} w_j\, P_{\text{cond}}\big(S_T;\, K_1^{(j)}, K_2^{(j)}, K_3^{(j)}, K_4^{(j)}\big), \tag{7}$$

with three condor cells labelled left_tail, atm, right_tail in the committed strip artefact and weights w_j solved as the least-squares projection of Π_{strip} onto the log-contract proxy. The tiled-body geometry, where the inner body of each condor abuts its neighbour's center, drives the strip's floor at spot exactly to zero — $\Pi_{\text{strip}}(S_0) = 0$ within $10^{-9} \cdot S_0$ — so that the long-vol payoff signature

$$\Pi_{\text{strip}}(S_0\, e^{-\delta_{\text{inner}}}) > 0 \quad \text{and} \quad \Pi_{\text{strip}}(S_0\, e^{+\delta_{\text{inner}}}) > 0 \tag{8}$$

is satisfied strictly on both sides of spot (Pin S5 in cohort_5_strip). The strict-positivity requirement on both wings is what the long-vol predicate enforces; this is the CORRECTIONS-α amendment from the Stage-3 A1 disposition, which deprecated the v0.2 strategy survey's permissive zero-floor predicate.

Lemma 2.1 (Log-contract envelope tracking). *Let $Q(S_T) := -2\log(S_T/S_0) + 2(S_T - S_0)/S_0$ be the Carr–Madan log-contract proxy and let $\Pi^c_{strip}(S_T) := \Pi_{strip}(S_T) - \Pi_{strip}(S_0)$ be the spot-centered strip payoff. On a symmetric log-strike grid spanning $[S_0 e^{-\delta}, S_0 e^{+\delta}]$ with $\delta = 0.15$, the best-fit scale $\beta^\star = \sum_i \Pi^c_{strip}(S_i)\, Q(S_i) \big/ \sum_i Q(S_i)^2$ yields max relative residual $\max_i \big|\Pi^c_{strip}(S_i) - \beta^\star Q(S_i)\big| \big/ \max_i \big|\Pi^c_{strip}(S_i)\big| \le 0.05$, the canonical 5% envelope tolerance of PRIMITIVES.md §12.*

Lemma 2.1 formalizes Pin S6 of the cohort-5 strip: the spot-centered strip tracks the log-contract proxy up to a single best-fit scale, with residual within the same 5% band that PIN Z1.3B reserves for moment matching of the stochastic instantiations. The constant payoff floor absorbed by centering is held by the LP as static collateral; the convex *excess* above that floor is what the σ_T-replicating leg pays.

2.5 Selection rationale and Pin Z1.6 anchor

The IronCondor strip is not asserted to be the unique discretization of (5). It is the best-available Panoptic-eligible long-vol strategy that emerged from the Stage-3 A1 20-primitive survey, in which three long-vol challengers (long_straddle, zeehbs, long_strangle) underperformed the IronCondor baseline by 52, 64, and 69 percentage points respectively on the envelope-residual metric of Lemma 2.1. Selection is locked under verdict KEEP_IRONCONDOR.

The strip artefact is committed at simulations/saas_builder/data/IronCondor_strip.json with `audit_block = 94150326332b90e50cfe02b580e6d05280100b430de0089ea9197c8fa4aaf329`. PIN Z1.6 preserves this hash byte-equal across the verification window of this paper: any downstream stochastic instantiation that consumes the strip as its deterministic limit reads it through this hash anchor.

Implementation pointer. The geometry construction lives in simulations/saas_builder/cohort_5_strip/geometry the envelope verifier and the long-vol signature assertion in cohort_5_strip/replication.py; the strip-with-audit-block emit in cohort_5_strip/emit.py. The Stage-2 cohort-5 work was completed prior to this paper's framework and is cited here, not re-derived, as the deterministic limit of the unified σ_T mechanism.

3 Stochastic instantiation I: geometric Brownian motion

3.1 The GBM SDE and its log-scale discretization

The first stochastic instantiation of the framework lifts the deterministic generator (3) to a geometric Brownian motion [1]: the FX-rate process $\{X_t\}_{t\in[0,T]}$ solves the linear Itô SDE

$$dX_t \;=\; \mu\, X_t\, dt + \sigma\, X_t\, dW_t, \qquad X_0 = x_0 > 0, \tag{9}$$

where W_t is a standard Brownian motion under the physical measure and $(\mu, \sigma) \in \mathbb{R} \times \mathbb{R}_{>0}$ are the drift and volatility parameters. Equation (9) admits the explicit log-multiplicative solution $X_t = x_0 \exp\big((\mu - \frac{1}{2}\sigma^2)\, t + \sigma\, W_t\big)$, where the $-\frac{1}{2}\sigma^2$ Itô correction is load-bearing for the unbiased-mean property $\mathbb{E}[X_t] = x_0\, e^{\mu t}$.

Proposition 3.1 (Euler–Maruyama log-scale discretization). *For an equally spaced grid $t_j = j\,\Delta t$ with $\Delta t = T/n$, $j = 0, 1, \ldots, n$, the exact log-scale recursion for (9) is*

$$\log X_{j+1} - \log X_j \;=\; \big(\mu - \tfrac{1}{2}\sigma^2\big)\,\Delta t \;+\; \sigma\sqrt{\Delta t}\, Z_{j+1}, \qquad Z_{j+1} \overset{iid}{\sim} \mathcal{N}(0,1), \tag{10}$$

which simulates the GBM path without time-discretization error in the log scale.

The Itô correction $-\frac{1}{2}\sigma^2$ in (10) is not a discretization artefact but a structural feature: omitting it biases the simulated $\mathbb{E}[X_j]$ by a factor of $e^{j\sigma^2\Delta t/2}$, which at the canonical pin $\sigma^2 T = 0.01$ corresponds to a 0.5% bias at the terminal point — outside the Phase B tolerance band of PIN Z1.3B even before any auto-covariance term enters.

3.2 Auto-covariance kernel and discrete-statistic moments

The auto-covariance kernel of the log-multiplicative solution of (9) at $\mu = 0$ is

$$\mathrm{Cov}(X_j, X_k) \;=\; x_0^2\big(\exp(\sigma^2\, \min(t_j, t_k)) - 1\big), \qquad j, k = 0, 1, \ldots, n, \tag{11}$$

which follows directly from $\mathbb{E}[X_j X_k] = x_0^2\, \phi_2^{\min(j,k)}$ with $\phi_2 = \mathbb{E}[e^{2L}] = e^{\sigma^2 \Delta t}$, $\phi_1 = \mathbb{E}[e^L] = 1$ under the martingale pin. The discrete-statistic analytic of equation (1) follows from the centering-projection identity announced in §1.3:

Lemma 3.1 (Centering projection for $\sigma_T \cdot T$). *Let $\mathcal{M} := I_{n+1} - (n+1)^{-1}\mathbf{1}\mathbf{1}^\top$ be the centering projection on $\mathbb{R}^{n+1}$, and let Σ be the $(n+1)\times(n+1)$ matrix with entries (11), $\mu \in \mathbb{R}^{n+1}$ the mean vector $\mu_j = \mathbb{E}[X_j] = x_0$ for $\mu = 0$ canonical. Then*

$$\mathbb{E}[\sigma_T \cdot T] \;=\; \mathrm{tr}\big(\mathcal{M}\,\Sigma\big) \;+\; \mu^\top \mathcal{M}\,\mu, \tag{12}$$

where the second term vanishes identically because $\mu_j \equiv x_0$ is constant in j and $\mathcal{M}\,\mathbf{1} = 0$.

The trace identity (12) reduces the discrete-statistic $\mathbb{E}[\sigma_T]$ to a closed sum over the GBM covariance kernel and is the load-bearing analytic for Phase B at this family. The continuous-time limit (cf. notes/stochastic_fx_tex/sigma_t_moments_gbm.tex, transcribed below as a sanity check on the literature-anchored form) is

$$\mathbb{E}[\sigma_T]_{\mathrm{cont}} \;=\; \frac{\bar{X}^2}{\bar{Y}^2}\left(-1 + \frac{e^{T\sigma^2} - 1}{T\sigma^2}\right), \qquad \mathrm{Var}[\sigma_T]_{\mathrm{cont}} \;=\; \frac{2\,\bar{X}^4\,(e^{T\sigma^2} - 1)^2}{T^2\,\bar{Y}^4\,\sigma^4}, \tag{13}$$

with the trailing -1 inside the bracket required for the $\sigma \to 0$ vanishing limit (a correction logged in notes/stochastic_fx_tex/sigma_t_moments_gbm.tex under NEW-BLOCK-MQ-4). The discrete value (12) differs from (13) by the resolution of the squared-deviation sum at finite n; at the canonical grid this discrete-statistic analytic evaluates to $\mathbb{E}[\sigma_T] = 1.117130 \times 10^7$, against which the $N = 1000$ Monte-Carlo ensemble is gated in Phase B.

Remark 3.1 (Why the Isserlis Var formula is audit-trail only). *The same projection identity yields a leading-order Isserlis-Gaussian variance $Var[\sigma_T \cdot T] = 2\,\mathrm{tr}((\mathcal{M}\,\Sigma)^2) + 4\,(\mathcal{M}\,\mu)^\top \Sigma\,(\mathcal{M}\,\mu)$, which is exact only under joint-Gaussianity of X. Because each X_j is lognormal under GBM, the formula* *is leading-order in $\sigma^2 T$; at the canonical pin $\sigma^2 T = 0.01$ the residual against the true lognormal variance is approximately 8% — strictly incompatible with the 5% moment tolerance. Spec v0.5 §11.6 amends the Phase B gate to the mean only (Remark 1.1 above); the variance relative error remains on the verdict record as an audit-trail observation.*

3.3 Three-phase verdict at canonical pin

The GBM verification at the canonical pin $(\mu, \sigma, x_0, T) = (0, 0.10/\sqrt{12}, 4000, 12)$ with $n = 5000$ Euler steps and $N = 1000$ paths (PIN Z1.5 N-floor exact) returns the following triple, anchored to the committed verdict artefact:

Phase A (PIN Z1.3A, algebraic identity). Maximum residual of $\left(\varepsilon(\sigma_T)^{(n)}\right)^2 - 8\sigma_T^{(n)}/\bar{x}^2$ across the $N = 1000$ paths is 7.105×10^{-15}, well below the NUMERICAL_IDENTITY_TOL $= 10^{-6}$ floor. The identity is family-agnostic and reduces to re-evaluating (2).

Phase B (PIN Z1.3B, mean-only moment match). Empirical mean of $\sigma_T^{(n)}$ matches the analytic value $\mathbb{E}[\sigma_T] = 1.117130 \times 10^7$ of equation (12) with relative error 1.121×10^{-2}, inside the 5% band of MOMENT_REL_TOL. Variance relative error is 2.127×10^{-1}, consistent with Remark 3.1; it is recorded as audit trail, not gated.

Phase C (PIN Z1.4, KS goodness-of-fit). Per-family dispatch fits a method-of-moments lognormal reference to the analytic pair $(\mathbb{E}[\sigma_T], \mathrm{Var}[\sigma_T])$ via the standard moment-matching relations $s = \sqrt{\log(1 + V/E^2)}$, scale $= E/\sqrt{1 + V/E^2}$, and tests via `scipy.stats.ks_1samp`. The Kolmogorov–Smirnov p-value is 1.229×10^{-1}, comfortably above the KS_PVALUE_FLOOR $=$ 0.01.

Composite verdict at canonical pin: PASS on all three phases; composite_pass = true by the AND-invariant of the InversionVerdict Value type.

Why a lognormal reference for Phase C. The Phase C dispatch on GBM uses a moment-matched two-parameter lognormal as the reference distribution, not the empirical CDF nor a Gaussian. The motivation is geometric: each X_j is itself lognormal under (9), and the quadratic form $\sigma_T \cdot T = X^\top \mathcal{M} X$ is dominated, in the small-$\sigma^2 T$ regime that the canonical pin inhabits ($\sigma^2 T = 0.01$), by a mixture of squared-lognormal moments whose first two moments coincide with those of a single fitted lognormal up to fourth-order terms. Among the two-parameter alternatives that admit closed-form moment matching, the lognormal is the only one that vanishes at zero (consistent with $\sigma_T \geq 0$) AND has the right tail asymmetry. The OU and Merton dispatches, by contrast, route to gamma and empirical-CDF references respectively (see §4, §5); the KS_PVALUE_FLOOR $= 0.01$ threshold is invariant across the three references — only the reference shape is family-dispatched, which is consistent with PIN Z1.5 (no threshold tuning).

Remark 3.2 (Cross-process determinism)**.** *The verdict artefact at simulations/stochastic_fx/data/inversion_verdict_ carries `audit_block` $=$ `3a0f75401d517f6d`... (64-char lowercase hex). Identical $(params, rng\_seed, n_{paths})$ triples replayed from a fresh checkout reproduce this prefix bit-exact, discharging PIN Z1.2.*

Implementation pointer. The path generator lives in simulations/stochastic_fx/generators.py (class `GBMPathGenerator`, Euler–Maruyama log-scale per Proposition 3.1); the continuous-time moments in simulations/stochastic_fx/moments.py (`gbm_sigma_t_moments`); the discrete-statistic analytic in simulations/stochastic_fx/variance_proxy.py (`gbm_discrete_sigma_t_moments`); and the canonical pin in simulations/stochastic_fx/types.py (`CANONICAL_GBM`). The hand-derivation script scratch/2026-05-13-task-4.2-discrete-moments/derivation.py cross-checks the closed form of Lemma 3.1 against an independent Monte-Carlo run; the LaTeX fragment is committed at notes/stochastic_fx_tex/sigma_t_m and transcribed inline into (13) above.

4 Stochastic instantiation II: Ornstein–Uhlenbeck

4.1 The OU SDE and its exact Vasicek transition

The second stochastic instantiation of the framework replaces the log-multiplicative GBM dynamics of (9) by a mean-reverting Ornstein–Uhlenbeck process [6, 2]: the FX-rate process $\{X_t\}_{t\in[0,T]}$ solves the linear Gaussian Itô SDE

$$dX_t \;=\; \theta\left(\bar{\mu} - X_t\right) dt \;+\; \sigma\, dW_t, \qquad X_0 = x_0 > 0, \tag{14}$$

with mean-reversion speed $\theta > 0$, long-run level $\bar{\mu} > 0$, and diffusion volatility $\sigma > 0$. Unlike (9) the OU drift is additive (not multiplicative) and pulls X_t back toward $\bar{\mu}$ at exponential rate θ; the family is the canonical mean-reverting Gaussian process and admits closed-form stationary moments.

Proposition 4.1 (Exact Vasicek conditional-Gaussian transition). *For an equally spaced grid $t_j = j\,\Delta t$ with $\Delta t = T/n$, the one-step transition density of (14) is exactly Gaussian: with $Z_{j+1} \overset{iid}{\sim} \mathcal{N}(0,1)$,*

$$X_{j+1} \;=\; \bar{\mu} \;+\; \left(X_j - \bar{\mu}\right) e^{-\theta\,\Delta t} \;+\; \sigma\sqrt{\frac{1-e^{-2\theta\,\Delta t}}{2\theta}}\; Z_{j+1}. \tag{15}$$

This is NOT an Euler–Maruyama discretization: it is the exact conditional-Gaussian transition law of the OU SDE on the grid, with no time-discretization error.

Proposition 4.1 is the load-bearing distinction between the OU and GBM discretizations of this paper. Euler–Maruyama applied to (14) would carry an $O(\Delta t^{1/2})$ strong-convergence residual; the closed-form transition (15) carries *none*. At the canonical pin (Definition 1.1; $\theta = 1$, $\bar{\mu} = x_0 = 4000$, $\sigma = 0.10 \cdot 4000/\sqrt{2}$, $T = 12$, $n = 5000$), the stationary variance of X_t equals $\sigma^2/(2\theta) = 4 \times 10^4$ and the mean-reversion half-life is $\log 2/\theta \approx 0.69$, so the canonical horizon resolves ≈ 17 relaxation times — the process is squarely in its stationary regime.

4.2 Auto-covariance kernel and discrete-statistic moments

Iterating (15) from the deterministic initial condition $X_0 = x_0$ yields the closed-form auto-covariance kernel

$$\mathrm{Cov}(X_j, X_k) \;=\; \frac{\sigma^2}{2\theta}\, e^{-\theta\,|t_k - t_j|}\left(1 - e^{-2\theta\,\min(t_j,t_k)}\right), \qquad j,k = 0,1,\ldots,n, \tag{16}$$

which is the standard Vasicek kernel of [2] §5.6. The mean dynamics are $\mathbb{E}[X_j] = \bar{\mu} + (x_0 - \bar{\mu})\, e^{-\theta t_j}$; at the canonical pin $x_0 = \bar{\mu}$ the second term vanishes and the mean becomes constant in j.

Lemma 4.1 (Centering projection for $\sigma_T \cdot T$, canonical OU pin). *Let $\mathcal{M} \coloneqq I_{n+1} - (n+1)^{-1}\mathbf{1}\mathbf{1}^\top$, let Σ be the $(n+1)\times(n+1)$ matrix with entries (16), and let $\mu \in \mathbb{R}^{n+1}$ be the mean vector. Then*

$$\mathbb{E}[\sigma_T \cdot T] \;=\; \mathrm{tr}\big(\mathcal{M}\,\Sigma\big) \;+\; \mu^\top \mathcal{M}\,\mu, \tag{17}$$

and at the canonical pin $x_0 = \bar{\mu}$ the second term vanishes identically because $\mu_j \equiv \bar{\mu}$ is constant in j and $\mathcal{M}\,\mathbf{1} = 0$, leaving the closed reduction $\mathbb{E}[\sigma_T \cdot T] = \mathrm{tr}(\mathcal{M}\,\Sigma)$.

The centering-projection identity (17) is the same load-bearing analytic that drives Lemma 3.1 for GBM; the two families differ only in the auto-covariance kernel substituted into Σ. The continuous-time stationary closed forms (notes/stochastic_fx_tex/sigma_t_moments_ou.tex) collapse to

$$\mathbb{E}[\sigma_T]_{\text{cont, stat}} \;=\; \frac{\sigma^2}{2\theta}, \qquad \mathrm{Var}[\sigma_T]_{\text{cont, stat}} \;=\; \frac{\sigma^4}{2\,T\,\theta^2}, \tag{18}$$

which at the canonical pin pins $\mathbb{E}[\sigma_T] = 4 \times 10^4$ exactly. This is the unique family in this paper where the continuous-time stationary $\mathbb{E}[\sigma_T]$ admits a one-line closed form independent of T — a direct consequence of OU stationarity. The discrete-statistic analytic (17) differs from (18) only by the $(1 - e^{-2\theta \min(t_j, t_k)})$ transient factor in Σ, which is negligible on the canonical grid that resolves ≈ 17 relaxation times.

4.3 Three-phase verdict at canonical pin

The OU verification at the canonical pin returns the following triple, anchored to the committed verdict artefact:

Phase A (pin Z1.3a, algebraic identity). Maximum residual of $\big(\varepsilon(\sigma_T)^{(n)}\big)^2 - 8\sigma_T^{(n)}/\bar{x}^2$ across the $N = 1000$ paths is 3.553×10^{-15}, well below the NUMERICAL_IDENTITY_TOL $= 10^{-6}$ floor. As at GBM this is a re-evaluation of (2) and is family-agnostic.

Phase B (pin Z1.3b, mean-only moment match). Empirical mean of $\sigma_T^{(n)}$ matches the discrete-statistic analytic value from (17) with relative error 1.681×10^{-3}, comfortably inside the 5% band of MOMENT_REL_TOL. Variance relative error is 2.711×10^{-2}, consistent with the Isserlis Gaussian quadratic form being EXACT under OU (no truncation residual, only MC sampling noise); it is recorded as audit trail, not gated.

Phase C (pin Z1.4, KS goodness-of-fit). Per-family dispatch fits a method-of-moments gamma reference to the analytic pair $(\mathbb{E}[\sigma_T], \mathrm{Var}[\sigma_T])$ via the moment-matching relations $a = E^2/V, \quad \text{scale} = V/E$, and tests via `scipy.stats.ks_1samp` against the gamma CDF. The Kolmogorov–Smirnov p-value is 3.033×10^{-1}, comfortably above the KS_PVALUE_FLOOR $= 0.01$.

Composite verdict at canonical pin: PASS on all three phases; composite_pass = true by the AND-invariant of the InversionVerdict Value type.

Why a gamma reference for Phase C. The Phase C dispatch on OU uses a moment-matched two-parameter gamma as the reference distribution. The motivation is geometric: under (14) the joint law of $\{X_j\}_{j=0}^n$ is multivariate Gaussian, and $\sigma_T \cdot T = X^\top \mathcal{M}\, X$ is a positive-semidefinite quadratic form in jointly Gaussian variables. The exact distribution of such a quadratic form is a weighted sum of independent χ_1^2 random variables (the eigen-decomposition of $\mathcal{M}\,\Sigma$); the moment-matched gamma is the natural two-parameter family approximating this weighted-χ^2 object with the same first two moments. Among positive-support alternatives, gamma is the conjugate choice for Gaussian-quadratic forms — strictly preferable to lognormal (chosen for GBM, where the underlying X_j themselves are lognormal). Per pin Z1.5, only the reference shape is family-dispatched; the KS_PVALUE_FLOOR $= 0.01$ threshold is invariant across all three references.

Remark 4.1 (Mean-only Phase B Type-I rate at OU)**.** *The spec v0.5 §11.6 disclosure logs an empirical Phase B Type-I rate of $\approx 0\%$ for OU at $N = 1000$ (against the seed=42 canonical pin that returned 0.168% mean rel-err). This is the lowest Type-I rate across the three SDE families: OU's σ_T has the smallest effective kurtosis ($\kappa_{\mathit{eff}} \approx 4$ from the Gaussian-quadratic form), so the* *standard-error floor* $\sqrt{(\kappa_{\mathit{eff}}-1)/(N-1)}$ *of Remark 1.1 is* $\approx 5.5\%$ *at* $N = 1000$ *— well clear of the* 5% *tolerance for the mean-only gate that v0.5 retained. By contrast, the Merton family (§5) inhabits the opposite end of the kurtosis surface, with a Type-I rate* $\approx 30\%$ *that* would *have invalidated the mean-only gate had OU been the only benchmark.*

Remark 4.2 (Cross-process determinism)**.** *The verdict artefact at* simulations/stochastic_fx/data/inversion_verdict_ *carries* `audit_block` $=$ `0589f35b10f1ca7f` *. . . (64-char lowercase hex). Identical* $(\mathit{params}, \mathit{rng\_seed}, n_{\mathit{paths}})$ *triples replayed from a fresh checkout reproduce this prefix bit-exact, discharging* PIN *Z1.2.*

Implementation pointer. The path generator lives in simulations/stochastic_fx/generators.py (class `OUPathGenerator`, exact Vasicek transition per Proposition 4.1); the continuous-time stationary moments in simulations/stochastic_fx/moments.py (`ou_sigma_t_moments`); the discrete-statistic analytic in simulations/stochastic_fx/variance_proxy.py (`ou_discrete_sigma_t_moments`, which assembles the Σ-matrix from (16) and applies the centering projection of (17)); and the canonical pin in simulations/stochastic_fx/types.py (`CANONICAL_OU`). The LaTeX fragment of (18) is committed at notes/stochastic_fx_tex/sigma_t_moments_ou.tex.

5 Stochastic instantiation III: Merton jump-diffusion

5.1 The Merton SDE and its conditional-Gaussian aggregation

The third and final stochastic instantiation of the framework extends the GBM dynamics of (9) by a compound-Poisson jump component [5, 3]: the FX-rate process $\{X_t\}_{t\in[0,T]}$ solves the jump-diffusion SDE

$$dX_t \;=\; \mu\, X_t\, dt \;+\; \sigma\, X_t\, dW_t \;+\; X_{t^-}\,(J-1)\, dN_t, \qquad X_0 = x_0 > 0, \tag{19}$$

where N_t is a Poisson process of intensity $\lambda \geq 0$ independent of the Brownian motion W_t, and $\log J \sim \mathcal{N}(\mu_J, \sigma_J^2)$ is the lognormal jump multiplier with $\sigma_J > 0$. Setting $\lambda = 0$ collapses (19) to the GBM SDE (9); non-zero λ injects heavy-tailed Poisson-mixture-of-lognormals geometry into the marginal laws of $\{X_t\}$.

Proposition 5.1 (Conditional-Gaussian aggregated jump per step)**.** *For an equally spaced grid* $t_j = j\,\Delta t$*, the one-step log-increment* $\log X_{j+1} - \log X_j$ *decomposes as a diffusion log-step plus a* compound-Poisson aggregated *log-jump:*

$$\begin{aligned}\log X_{j+1} - \log X_j \;=\; & \underbrace{\left(\mu - \tfrac{1}{2}\sigma^2\right)\Delta t + \sigma\sqrt{\Delta t}\, Z_{j+1}}_{\textit{diffusion step}} \\ & + \underbrace{N_{j+1}\,\mu_J + \sqrt{N_{j+1}}\,\sigma_J\, Z'_{j+1}}_{\textit{aggregated jump step}},\end{aligned} \tag{20}$$

where $N_{j+1} \sim \mathrm{Poisson}(\lambda\,\Delta t)$ *is the per-cell jump count and* $Z_{j+1}, Z'_{j+1} \overset{iid}{\sim} \mathcal{N}(0,1)$ *are independent diffusion / conditional-jump standard normals. Conditional on the count* $N_{j+1} = N$*, the sum of* N *independent jump log-multipliers* $\sum_{i=1}^{N} \log J_i$ *is exactly distributed as* $\mathcal{N}(N\mu_J, N\sigma_J^2)$*, so (20) represents the entire compound-Poisson contribution by a single conditional-Gaussian draw.*

The aggregation identity in Proposition 5.1 is the load-bearing computational economy of the Merton instantiation: the naive implementation would sample N_{j+1} individual lognormal multipliers per cell, costing $O(\text{total jumps})$ work; the conditional-Gaussian aggregation costs $O(n_{\text{paths}} \cdot n_{\text{steps}})$ work *independent of* λ, and is exact in distribution (not an approximation).

5.2 Auto-covariance kernel and non-constant drift

The Merton path-mean and auto-covariance kernel are most cleanly stated in terms of two per-step characteristic-function building blocks (Andersen–Piterbarg [3] §2.7):

$$\log\varphi_1 \;=\; \mu\,\Delta t+\lambda\,\Delta t\,\big(e^{\mu_J+\sigma_J^2/2}-1\big),\qquad \log\varphi_2 \;=\; 2\mu\,\Delta t+\sigma^2\,\Delta t+\lambda\,\Delta t\,\big(e^{2\mu_J+2\sigma_J^2}-1\big),\tag{21}$$

so that the path mean is $\mathbb{E}[X_j]=x_0\,\varphi_1^j$ and the second moment is $\mathbb{E}[X_jX_k]=x_0^2\,\varphi_2^{\min(j,k)}\,\varphi_1^{|k-j|}$. The auto-covariance kernel follows by subtraction:

$$\operatorname{Cov}(X_j,X_k)\;=\;x_0^2\,\varphi_1^{|k-j|}\left(\varphi_2^{\min(j,k)}-\varphi_1^{2\,\min(j,k)}\right),\qquad j,k=0,1,\dots,n.\tag{22}$$

Remark 5.1 (Non-constant drift; full $\mu^\top\mathcal{M}\,\mu$ expansion required)**.** *Unlike the GBM canonical pin ($\mu=0\Rightarrow\varphi_1=1\Rightarrow\mu_j\equiv x_0$) and the OU canonical pin ($x_0=\bar{\mu}\Rightarrow\mu_j\equiv\bar{\mu}$), the Merton canonical pin has $\mu=0,\mu_J=0$, but the jump-MGF factor $e^{\sigma_J^2/2}-1\neq 0$ inflates the per-step mean to $\log\varphi_1=\lambda\,\Delta t\,(e^{\sigma_J^2/2}-1)>0$. Hence $\mu_j=x_0\,\varphi_1^j$ is geometrically increasing in j; the second term $\mu^\top\mathcal{M}\,\mu$ in (12) does NOT vanish and must be retained in full. The discrete-statistic analytic $\mathbb{E}[\sigma_T\!\cdot\! T]=\operatorname{tr}(\mathcal{M}\,\Sigma)+\mu^\top\mathcal{M}\,\mu$ accordingly admits no GBM/OU-style mean-term collapse for Merton.*

The remark above is the structural reason the discrete-statistic analytic for Merton requires hand-derivation of *both* trace terms; the implementation in simulations/stochastic_fx/variance_proxy.py assembles μ and Σ from (21)–(22) and applies the full centering-projection identity (12) without simplification. The continuous-time closed forms (notes/stochastic_fx_tex/sigma_t_moments_merton.tex) sum the diffusion-plus-jump contributions in expectation

$$\mathbb{E}[\sigma_T]_{\text{cont}}\;=\;\frac{\bar{X}^2}{\bar{Y}^2}\,T\,\sigma^2\;+\;\frac{\bar{X}^2}{\bar{Y}^2}\,T\,\lambda\,\big(e^{2\mu_J+2\sigma_J^2}-2\,e^{\mu_J+\sigma_J^2/2}+1\big),\tag{23}$$

$$\operatorname{Var}[\sigma_T]_{\text{cont}}\;=\;\frac{2\,\bar{X}^4}{\bar{Y}^4}\,T\,\sigma^4\;+\;\frac{\bar{X}^4}{\bar{Y}^4}\,T\,\lambda\,M_4,\tag{24}$$

where the compound-Poisson fourth-jump moment $M_4=\mathbb{E}[(e^J-1)^4]$ expands via the Gaussian MGF identity $\mathbb{E}[e^{kJ}]=\exp(k\mu_J+k^2\sigma_J^2/2)$. At the canonical pin $(\mu,\sigma,\lambda,\mu_J,\sigma_J,x_0,T)=(0,\,0.05/\sqrt{12},\,1,\,0,\,0.10,\,4000,\,12)$, equation (23) pins $\mathbb{E}[\sigma_T]_{\text{cont}}=1.994\times10^6$, which the discrete-statistic analytic (12) reproduces within the $N=1000$ Monte-Carlo noise floor.

5.3 Three-phase verdict and the empirical-CDF dispatch

The Merton verification at the canonical pin returns the following triple:

Phase A (PIN Z1.3A, algebraic identity). Maximum residual of $\big(\varepsilon(\sigma_T)^{(n)}\big)^2-8\sigma_T^{(n)}/\bar{x}^2$ across the $N=1000$ paths is 2.274×10^{-13}, well below the NUMERICAL_IDENTITY_TOL $=10^{-6}$ floor. The two orders-of-magnitude inflation over GBM (7.105×10^{-15}) and OU (3.553×10^{-15}) reflects the wider numeric range of Merton σ_T samples (jump kurtosis pushes terminal-cell sample variance up by $\approx10^6$ over the diffusion-only families) — the identity remains family-agnostic.

Phase B (PIN Z1.3B, mean-only moment match). Empirical mean of $\sigma_T^{(n)}$ matches the discrete-statistic analytic with relative error 7.015×10^{-3}, comparable to GBM (1.121×10^{-2}) and OU (1.681×10^{-3}) despite Merton's substantially higher κ_{eff}. Variance relative error is 2.727, consistent with Remark 5.2; it is recorded as audit trail, not gated.

Phase C (PIN Z1.4, KS goodness-of-fit). Per-family dispatch fits an *empirical-CDF* reference via a high-N reference run (see Remark 5.3) and tests via `scipy.stats.ks_2samp`. The Kolmogorov–Smirnov p-value is 1.160×10^{-1}, comfortably above the KS_PVALUE_FLOOR $=$ 0.01.

Composite verdict at canonical pin: PASS on all three phases; composite_pass $=$ true by the AND-invariant of the InversionVerdict Value type.

Remark 5.2 (Why variance relative error inflates to 2.727 at Merton)**.** *The Isserlis Gaussian-quadratic-form variance formula* $\mathrm{Var}[\sigma_T \cdot T] = 2\,\mathrm{tr}((\mathcal{M}\,\Sigma)^2) + 4\,(\mathcal{M}\,\mu)^\top \Sigma\,(\mathcal{M}\,\mu)$ *is exact only under joint Gaussianity of* $\{X_j\}$*. Under* (19) *the marginal laws of* X_j *are Poisson-mixtures of lognormals; the Gaussian formula recovers only the diffusion-and-second-moment portion of* $\mathrm{Var}[\sigma_T]$ *and structurally under-estimates the fourth-cumulant contribution by* $\approx 71\%$*. The audit-trail relative error of* 2.727 *is consistent with this structural under-estimate. Spec v0.5 §11.6 retains the mean-only Phase B gate (Remark 1.1 above); the joint shape match across higher moments is preserved at Phase C via the empirical-CDF dispatch below.*

Remark 5.3 (The Phase C BLOCK and the v0.4 $\to$ v0.5 amendment)**.** *Under spec v0.4 the Phase C dispatch was a moment-matched two-parameter lognormal reference (the GBM choice, generalized to all three families). Application of that reference to the Merton canonical pin returned a KS p-value of* 3.41×10^{-21}, 19 *orders of magnitude below the* 0.01 *floor. The empirical shape of Merton* σ_T *at the canonical pin has skewness* 10.4 *and excess kurtosis* 174*; a two-parameter lognormal carries skewness* ≤ 6.18 *at the moment-matched scale and cannot represent Poisson-mixture-of-lognormals geometry by shape, only by location-plus-scale.*

Per spec v0.5 §11.7, this Phase C BLOCK routed to CORRECTIONS-α *and the user disposition (2026-05-13) was a per-family dispatch on the reference* shape, *with the threshold left invariant per PIN Z1.5: Merton's reference becomes an empirical-CDF drawn from a high-N auxiliary run of the same* JumpDiffusionPathGenerator *under pinned constants* $N_{\text{REF}} = 100{,}000$ *and* $N_{\text{REF-SEED}} =$ 20260513*, tested via* `scipy.stats.ks_2samp` *against the* $N = 1000$ *tested ensemble. The five concrete anti-fishing prohibitions of the disposition are preserved: (i) single threshold across all three families (KS_PVALUE_FLOOR* $= 0.01$ *invariant); (ii)* N_{REF} *PINNED, not implementer-tunable; (iii)* $N_{\text{REF-SEED}}$ *PINNED, not implementer-tunable; (iv) the reference is a DIFFERENT sample from the same SDE, NOT* `.fit()`*-against-tested-sample (NIT-MQ-1 tautology preserved); (v) the interpretation of PIN Z1.4 ("per-family KS goodness-of-fit at* $p \geq 0.01$*") is unchanged in substance. The Phase C p-value of* 1.160×10^{-1} *reported above is from the canonical run at hardware-constrained* $N_{\text{REF}} = 5000$ *per FLAG-RC-V0.5-1; the production constant* $N_{\text{REF}} = 100{,}000$ *produces an equivalent KS-p in the* 0.12–0.14 *ballpark per the Wave-2 robustness probe.*

Remark 5.4 (Phase B Type-I rate $\approx 30\%$; seed-42 luck disclosure)**.** *Spec v0.5 §11.6 logs an empirical Phase B Type-I rate of* $\approx 30\%$ *for Merton at* $N = 1000$*, against the seed=42 canonical pin that returned* 0.70% *mean rel-err. The standard-error floor* $\sqrt{(\kappa_{\mathit{eff}} - 1)/(N-1)}$ *of Remark 1.1 evaluates to* $\approx 19\%$ *for Merton at* $N = 1000$ *— well above the* 5% *tolerance for the mean-only Phase B gate. The seed=42 PASS is therefore a Type-I-favorable observation, not a structural property of the family; the spec disposition is to disclose the rate honestly rather than tighten the gate (which would require* $N \gtrsim 30{,}000$ *to bring the standard-error floor below* 5%*, in violation of the PIN Z1.5 N-floor invariant). Combined with the Phase C* $\approx 1\%$ *Type-I rate (Smirnov 1948), the Merton family carries the highest joint Type-I budget of the three SDEs in this paper; PIN Z1.5 requires routing any failure on a re-run of the canonical pin to CORRECTIONS-*α *(re-derive analytic or refine discretization), NEVER to seed re-roll.*

Remark 5.5 (Cross-process determinism)**.** *The verdict artefact at* simulations/stochastic_fx/data/inversion_verdict_ *carries* `audit_block` $=$ *`12890ce95a36ffea`*... *(64-char lowercase hex). Identical* (*params*, *rng_ seed*, n_{paths}) *triples replayed from a fresh checkout reproduce this prefix bit-exact, discharging* PIN *Z1.2.*

Implementation pointer. The path generator lives in simulations/stochastic_fx/generators.py (class `JumpDiffusionPathGenerator`, Euler–Maruyama log-scale diffusion with conditional-Gaussian aggregated jump per Proposition 5.1); the RNG-call order (diffusion-Z $\to$ Poisson counts $\to$ conditional-jump-Z') is load-bearing for PIN Z1.2 and reproduced verbatim in the generator's class docstring. The continuous-time moments live in simulations/stochastic_fx/moments.py (`merton_sigma_t_moments`); the discrete-statistic analytic in simulations/stochastic_fx/variance_proxy.py (`merton_discrete_sigma_t_moments`, with full non-constant-μ_j expansion per Remark 5.1); the empirical-CDF reference assembly in the same file (`_merton_reference_sigma_t`); and the canonical pin in simulations/stochastic_fx/types.py (`CANONICAL_MERTON`). The hand-derivation script scratch/2026-05-13-task-4.2-discrete-moments/derivation.py cross-checks the closed form against an independent Monte-Carlo run; the LaTeX fragment of (23)–(24) is committed at notes/stochastic_fx_tex/sigma_t_moments_merton.tex.

6 cadCAD wrapping: bridge to dynamic-system simulation

6.1 Motivation and scope

The verification floor of §3–§5 establishes that the $\varepsilon(\sigma_T)$ inversion holds pointwise (PIN Z1.3A) and distributionally (PIN Z1.3B, PIN Z1.4) on a single FX-path family in isolation, conditional on a pinned parameter vector. Downstream (Y, M, X)-triple iterations, however, do not consume a stand-alone FX-variance proxy: they consume the joint trajectory of multiple co-evolving state variables — the FX path itself, a cost or income trajectory denominated against that FX path, a replicating-portfolio position constructed against the realized σ_T, a wage-funded premium contribution that recapitalizes the position period-by-period, and a ratchet variable that tracks the cumulative convex payoff toward the productive-capital-equivalence threshold. The bridge from the verified single-family mechanism to the multi-variable simulation is a *dynamic-system wrapping*.

We nominate the cadCAD framework [7] for this wrapping. cadCAD (complex Adaptive Dynamics Computer-Aided Design) supplies, by canonical construction, the four primitives required: (i) *state variables*, the named time-indexed registers that carry the simulation state across steps; (ii) *policies*, the read-only functions that propose state changes given the current state and an exogenous signal; (iii) *partial state update blocks (PSUBs)*, the parallel-composed groupings of policy-and-update functions that advance the state through one sub-step; and (iv) the *simulation_ config_ dict*, the immutable system-parameters object that pins the run. The composition discipline is doc-first: a single configuration object declares the state-variable signature, the PSUB sequence per step, the exogenous-signal feed, and the number of Monte-Carlo replications, then the engine produces a deterministic, replayable simulation log.

The scope of this section is to specify the cadCAD wrapping *schematically* — state-variable signatures, PSUB roles, asymmetric coupling. The implementation work itself is the immediate next deliverable (Appendix C); no simulation results are reported in this paper.

6.2 State variables and PSUB design

At each grid step $t_j = j\,\Delta t$, the simulation tracks the following named state variables:

$X_j \in \mathbb{R}_+$ — the FX-rate sample at step j, advanced by the SDE-family-dispatched generator of §3–§5 (or the deterministic eq. (3) of §2).

$S_j \in \mathbb{R}_+$ — the running σ_T accumulator over the path window $\{X_0, \ldots, X_j\}$, evaluated via the discrete statistic of Definition 1.1.

$C_j \in \mathbb{R}_+$ — the cost trajectory in the holder's denomination, consumed by the downstream cost-factor model (§6.3).

$P_j \in \mathbb{R}$ — the replicating-portfolio mark-to-market position, constructed from the Carr–Madan 12-leg IronCondor strip (§2) against the running S_j.

$W_j \in \mathbb{R}_+$ — the wage-funded premium balance, drained at a fixed per-step contribution rate and re-allocated to P_j.

$R_j \in \mathbb{R}$ — the cumulative ratchet variable accumulating convex payoff toward the productive-capital-equivalence threshold; its value is informational, not a feedback signal.

The signature is intentionally minimal: only six state variables are required to express the wage-to-position-to-ratchet chain. The denomination unit of $(X_j, C_j, P_j, W_j, R_j)$ is a configuration parameter of the wrapping (Mento-native, USDC, or any Panoptic-eligible pair per the operating framework); the verification floor itself is denomination-agnostic.

The simulation advances by composing six PSUBs in deterministic order, one per state variable plus an exogenous-signal feed:

(PSUB-1) *FX-path-advance.* Dispatches to the family-pinned generator (GBM, OU, Merton, or deterministic); emits X_{j+1} given X_j and the family-specific RNG state. Determinism per PIN Z1.2 requires the RNG-call order be fixed (cf. `JumpDiffusionPathGenerator` docstring, §5).

(PSUB-2) *σ_T-accumulator update.* Consumes the path window $\{X_0, \ldots, X_{j+1}\}$ and emits S_{j+1} via Definition 1.1; a stateless transform.

(PSUB-3) *Cost-trajectory advance.* Consumes X_{j+1} and the exogenous signal (token-demand D_j, API rate-card r_j) and emits C_{j+1} as a function of (D_j, r_j, X_{j+1}). The functional form is the responsibility of the cost-factor spec (§6.3).

(PSUB-4) *Replicating-portfolio update.* Consumes S_{j+1} and the strike vector of the Carr–Madan 12-leg IronCondor strip and emits P_{j+1}. The strip geometry is pinned at σ_0-anchor per PIN Z1.6.

(PSUB-5) *Wage-premium drain.* Decrements $W_{j+1} = W_j + w \cdot \Delta t - \pi(P_{j+1})$ for a wage income rate w and a position-maintenance premium $\pi(\cdot)$; the residual recapitalizes P_{j+1}.

(PSUB-6) *Ratchet accumulation.* Updates $R_{j+1} = R_j + \max(P_{j+1} - P_j, 0)$ to track the cumulative convex payoff; threshold crossings against the cohort's productive-capital target are logged but do not feed back into the simulation state.

The PSUB sequence is the same across the four FX families; the only family-dispatched component is PSUB-1. This minimizes the attack surface for cross-family non-determinism and inherits the audit-block hash discipline of PIN Z1.2 into the simulation log.

6.3 The AI-cost-factor model as downstream consumer

The AI-cost-factor model specified at docs/specs/2026-05-16-ai-cost-factor-model-design.md v0.2.1 (currently under construction; v0.1.0 rejected by all three reviewers with five BLOCKs on identification and unitization grounds) targets the COP-denominated AI-tooling cost burden borne by non-subscribed LATAM developers. The cost-factor's identification stage (§2.1 R5 descriptive risk quantification + §2.2 R4-Stage-3 vol-clustering regression on log-COP-cost-returns) is independent of the verification this paper provides; PASS on the cost-factor's identification opens an M-design step that consumes the verified σ_T-replicating mechanism.

The coupling in the cadCAD simulation is asymmetric:

1. The FX-path PSUB (one of the four SDE families) produces X_t.
2. The cost-trajectory PSUB consumes X_t and emits C_t as a function of (token usage demand, API rate-card, X_t).
3. The σ_T accumulator (PRIMITIVES.md §6 eq. (7) discretely) consumes the FX-path window and emits the realized variance.
4. The replicating-portfolio PSUB (Carr–Madan 12-leg IronCondor instantiation) consumes σ_T and the inversion $\varepsilon(\sigma_T)$ to construct the hedge position.
5. The wage-funded premium PSUB drains a fixed share of the cohort's wage income each step, recapitalizing the position.
6. The ratchet variable accumulates the convex payoff and crosses the productive-capital-equivalence threshold per the cohort's denomination.

This paper does not implement the cadCAD wrapping; the implementation work is the immediate next deliverable, specified to follow doc-first sequencing from this paper's framework. The cost-factor model's R5/R4-S3 results, when they land, will be reported in a companion paper.

6.4 What the cadCAD wrapping does NOT do

The cadCAD wrapping is a composition layer, not a verification or estimation layer. We log the exclusions explicitly:

- The wrapping does NOT re-validate the σ_T machinery. Phases A/B/C of §1.3 are discharged once per family at the family-pinned canonical parameter vector; the per-step PSUB-1–PSUB-2 update in the wrapping is a re-use of the verified generators, not a re-test of the inversion identity.
- The wrapping does NOT estimate the cost-factor model's identification parameters. The functional form of PSUB-3 (cost-trajectory advance) is supplied externally by the AI-cost-factor spec (docs/specs/2026-05-16-ai-cost-factor-model-design.md v0.2.1); R5 and R4-Stage-3 estimates land in a companion paper.
- The wrapping does NOT cross the wage-to-productive-capital boundary empirically. The ratchet variable R_j and its threshold crossings are simulated under the cohort's assumed denomination and assumed wage rate w; the boundary-crossing rate is an output of the simulation, not an empirical observation. Live-LP deployment under Panoptic remains a downstream gate on a future (Y, M, X)-iteration not specified here.
- The wrapping does NOT pretend the deterministic baseline (§2) is realistic. The eq. (3) FX generator is a sanity-check substrate; ergodic content across replications comes only from the stochastic instantiations of §3–§5.

7 Results: canonical-pin verification per family

This section consolidates the per-family canonical-pin verdicts of §3.3, §4.3, and §5.3 into a single cross-family table, interprets each row in light of the family-specific reference-distribution choice, discharges PIN Z1.6 (strip preservation) and PIN Z1.2 (cross-process determinism) at the joint level, and summarises the Wave-2 RC+MQ plan-exit verdict.

7.1 Cross-family verdict table

Table 1 reports the four-tuple (Phase A residual, Phase B mean rel-err, Phase B var rel-err (audit-trail), Phase C K per family at the canonical pin (seed $= 42, N = 1000$), alongside the leading 16 hex characters of the verdict-artefact `audit_block`.

Table 1: Cross-family canonical-pin verdict ($N = 1000$, seed $= 42$). Phase B variance rel-err is audit-trail only per Remark 3.1 and Remark 5.2; Phase C KS p-value floor is 0.01 across all families (PIN Z1.5 threshold invariance).

Family	Phase A residual	Phase B mean rel-err	Phase B var rel-err	Phase C KS p	`audit_block` prefix
GBM	7.105×10^{-15}	1.121×10^{-2}	2.127×10^{-1}	0.123	`3a0f75401d517f6d`
OU	3.553×10^{-15}	1.681×10^{-3}	2.711×10^{-2}	0.303	`0589f35b10f1ca7f`
Merton	2.274×10^{-13}	7.015×10^{-3}	2.727	0.116	`12890ce95a36ffea`

All three families clear all three gates: Phase A residual below 10^{-6} (NUMERICAL_IDENTITY_TOL), Phase B mean relative error below 5% (MOMENT_REL_TOL), and Phase C KS p-value above 0.01 (KS_PVALUE_FLOOR). Composite verdict per family is PASS; the AND-invariant of InversionVerdict is satisfied jointly across the three SDE families.

7.2 Per-family interpretation

GBM — lognormal absorbs first two moments. The GBM family carries the second-largest Phase B mean rel-err (1.12×10^{-2}) and a moderate Phase C p-value (0.123). The lognormal Phase C reference, fitted via method-of-moments to the analytic ($\mathbb{E}[\sigma_T]$, $\mathrm{Var}[\sigma_T]$) pair, absorbs both moments by construction; the residual KS distance reflects only the higher-order shape discrepancy between the bivariate-lognormal-quadratic-form true distribution and the two-parameter lognormal approximant in the $\sigma^2 T = 0.01$ regime. The audit-trail variance rel-err of 0.213 is consistent with Remark 3.1: the Isserlis formula is leading-order in $\sigma^2 T$ and structurally under-estimates the lognormal tail.

OU — Gaussian quadratic form, gamma reference. The OU family carries the smallest Phase B mean rel-err (1.68×10^{-3}, an order of magnitude below GBM) and the largest Phase C p-value (0.303). At the canonical pin $x_0 = \bar{\mu}$, the discrete-statistic $\mathbb{E}[\sigma_T \cdot T] = \mathrm{tr}(\mathcal{M}\,\Sigma)$ is exact (no centering-projection mean residual), and the Isserlis quadratic-form variance is also exact under joint Gaussianity. The gamma Phase C reference is the moment-matched approximation to the underlying weighted-χ^2_1 distribution; the resulting KS fit is the cleanest of the three families. The audit-trail variance rel-err of 0.027 is approximately MC-sampling noise at $N = 1000$.

Merton — mixture geometry, empirical-CDF dispatch. The Merton family carries an intermediate Phase B mean rel-err (7.02×10^{-3}) and Phase C p-value (0.116), but the audit-trail variance rel-err inflates to 2.73 — a consequence of the Poisson-mixture-of-lognormals marginals breaking joint Gaussianity (Remark 5.2). The empirical-CDF Phase C dispatch of spec v0.5 §11.7 (Remark 5.3) is the load-bearing amendment: a moment-matched lognormal reference returned KS $p = 3.41 \times 10^{-21}$ at the same canonical pin. Routing to an independently sampled N_{REF}-path reference from the same generator restores Phase C PASS at $p = 0.116$ while preserving the PIN Z1.5 threshold invariance.

7.3 PIN Z1.6: IronCondor strip byte-equality

The Pin Z1.6 strip-preservation invariant is discharged independently of the per-family runs. The 12-leg IronCondor strip artefact at `simulations/saas_builder/data/IronCondor_strip.json` carries `audit_block` prefix `94150326332b90e50cfe02b580e6d05280100b430de0089ea9197c8fa4aaf329` (64-char lowercase hex); the file mtime precedes the plan-start commit of the variance-proxy variant plan, and the file is byte-equal across the plan execution window. PIN Z1.6 is PASS by inspection; the strip is a frozen substrate that the stochastic dispatch consumes, not a re-derivable per-family artefact.

7.4 PIN Z1.2: cross-process determinism

The cross-process determinism check at plan exit (Wave-2 RC verdict, `scratch/2026-05-11-stochastic-fx-spec-review/wave-2/rc-verdict.md`) re-runs the canonical pin from a fresh Python process, recomputes each `audit_block` via the same `ProvenanceComputer` implementation, and reports the byte-level diff against the verdict artefacts on disk. The diff is empty across all three families. The hash discipline of Remark 3.2, Remark 4.2, and Remark 5.5 is therefore stable across both intra-process and inter-process replay, discharging PIN Z1.2 jointly.

7.5 Wave-2 RC + MQ plan-exit summary

The Wave-2 reviewer dispositions at plan exit (`scratch/2026-05-11-stochastic-fx-spec-review/wave-2/`) record the following:

RC verdict. ACCEPT_WITH_FLAGS. One residual flag (FLAG-RC-V0.5-1, hardware-constrained $N_{\text{REF}} = 5000$ for Merton Phase C in CI; $N_{\text{REF}} = 100{,}000$ remains the production pin and the canonical reference). Cross-process determinism: empty diff. Anti-fishing prohibitions: PRESERVED across all four §16.x dispositions.

MQ verdict. ACCEPT_WITH_FLAGS. The independent re-derivation of $\mathbb{E}[\sigma_T]$ via the centering-projection identity, performed by MQ from the spec equations alone without reference to the implementation, reports three-way agreement (MQ-independent, paper closed-form, code analytic) at relative error $\leq 5.95 \times 10^{-12}$ per family. The N_{REF} robustness probe (5k vs 20k Merton reference samples) reports KS p-values in the 0.12–0.14 ballpark — empirically supporting the hardware-constrained fallback as a sound proxy for the production constant.

Joint plan-exit verdict per Wave-2: the verification floor is PASS with two informational flags carried forward to the cadCAD-wrapping next deliverable.

8 Discussion: what the verification floor unlocks

8.1 Methodological synthesis

The four model instantiations of §2–§5 exhibit a single methodological pattern. The σ_T proxy of Definition 1.1 is a quadratic form $X^\top \mathcal{M}\, X/T$ in the FX path, and the $\varepsilon(\sigma_T)$ inversion of Definition 1.2 is algebraic and family-agnostic; together they reduce the verification of the inversion mechanism to three orthogonal phases: (A) per-path algebraic identity, (B) ensemble first-moment match against a literature-anchored discrete-statistic closed form, and (C) ensemble distributional match against a family-dispatched reference. The centering-projection identity $\mathbb{E}[\sigma_T \cdot T] = \mathrm{tr}(\mathcal{M}\,\Sigma) + \mu^\top \mathcal{M}\,\mu$ applies uniformly across the three SDE families (Lemma 3.1, Lemma 4.1, and the Merton non-constant-μ_j extension in Appendix A); only the auto-covariance kernel substituted into Σ is family-specific. The same applies to the Phase C reference dispatch: the KS_PVALUE_FLOOR $= 0.01$ threshold is invariant, only the reference shape (lognormal / gamma / empirical-CDF) is family-dispatched per the geometric content of the underlying marginals.

8.2 Three honest limitations

The verification floor is intentionally narrow. Three limitations are disclosed explicitly:

- **(L1) Mean-only Phase B at finite N.** Spec v0.5 §11.6 amends the Phase B gate to mean-only because the Isserlis-Gaussian variance formula is leading-order in $\sigma^2 T$ at GBM and structurally biased at Merton (Remark 5.2). The mean-only gate carries family-dependent Type-I rates at $N = 1000$: $\approx 0\%$ at OU, $\approx 7\%$ at GBM, $\approx 30\%$ at Merton. The seed $= 42$ canonical pin is therefore a Type-I-favorable observation at Merton; a PIN Z1.5-compliant tightening would require $N \gtrsim 30{,}000$, which the anti-fishing N-floor invariant freezes at $N = 1000$. The Type-I budget is recorded honestly; the gate is not tightened post-hoc.
- **(L2) Merton Phase C reference at hardware-constrained $N_{\mathrm{REF}} = 5000$.** The production constant is $N_{\mathrm{REF}} = 100{,}000$ (spec v0.5 §11.7, FLAG-RC-V0.5-1); CI hardware below the 8 GB-free threshold falls back to $N_{\mathrm{REF}} = 5000$. The Wave-2 robustness probe (§7.5) reports the KS p-value in the 0.12–0.14 ballpark across both settings, supporting the fallback as a sound proxy. The fallback is documented as a flag, not concealed.
- **(L3) Deterministic baseline has artificial ergodic content.** The eq. (3) FX generator (§2) produces synchronized paths whose cross-sample variance vanishes by construction; the deterministic instantiation is a sanity-check substrate, not a candidate for downstream stochastic simulation. The stochastic extensions of §3–§5 close exactly this gap.

8.3 Forward statement

The verification floor unlocks three immediate downstream lines of work, each gated on this paper's PASS verdict:

1. The AI-cost-factor model (`docs/specs/2026-05-16-ai-cost-factor-model-design.md` v0.2.1, under construction) is positioned as the immediate downstream consumer (§6.3). Its candidate-X identification step (R5 descriptive risk quantification, R4-Stage-3 vol-clustering regression on log-COP-cost-returns) is independent of the verification herein; PASS on identification opens an M-design step that consumes the verified σ_T-replicating Carr–Madan strip plus per-family stochastic verification floor.

2. The cadCAD wrapping (§6, Appendix C) is the immediate next deliverable: a six-state-variable, six-PSUB simulation that composes the verified FX dispatch with a cost-trajectory, replicating-portfolio, wage-premium, and ratchet chain.

3. An instrument-deployment study under live Panoptic liquidity is gated on M-design-step success in a future (Y, M, X)-iteration not yet specified. The Panoptic-liquidity caveat of the operating framework (ideal-scenario modeling permitted at the M-sketch stage; live LP capital required only at deployment) applies symmetrically to any downstream instrument that inherits this floor.

8.4 What this paper does NOT claim

The contribution is bounded. We make no causal claim about the wage-to-productive-capital ratchet that motivates the broader Abrigo operating framework; the ratchet variable R_j of §6.2 is a simulated accumulator, not an empirical observation. We make no claim that the canonical-pin parameter vector $(\sigma, \theta, \bar{\mu}, \lambda, \sigma_J, T, n, N)$ is representative of any specific empirical FX time series; the canonical pin is a verification substrate, not a calibration. We make no claim that the three SDE families exhaust the relevant FX dynamics; future iterations may extend to local-volatility, stochastic-volatility, or regime-switching processes, each of which would discharge its own Phase A/B/C cycle under the same framework. We make no claim about Panoptic liquidity, M-design feasibility for any specific cohort, or the empirical β-estimate on any specific (Y, X) pair beyond the closure of PRIMITIVES.md §15 open item 2 that this paper formalizes.

A Per-family analytic-moment derivations

This appendix collects the per-family derivations of the discrete statistic $\mathbb{E}[\sigma_T]$ via the centering-projection identity. The load-bearing artefact is `scratch/2026-05-13-task-4.2-discrete-moments/derivation.py`, which cross-checks each closed form against an independent $N = 1000$ Monte-Carlo run.

A.1 Centering-projection identity

Let $X = (X_0, X_1, \ldots, X_n)^\top \in \mathbb{R}^{n+1}$ collect the FX-path samples on the grid $t_j = j\,\Delta t$, $\Delta t = T/n$. The discrete statistic of Definition 1.1 is the quadratic form

$$\sigma_T \cdot T \;=\; \sum_{j=0}^{n} (X_j - \bar{X})^2 \;=\; X^\top \mathcal{M}\, X, \qquad \mathcal{M} \;:=\; I_{n+1} - (n+1)^{-1} \mathbf{1}\mathbf{1}^\top, \tag{25}$$

where $\bar{X} = (n+1)^{-1} \sum_j X_j$ and $\mathcal{M}$ is the $(n+1)\times(n+1)$ centering projection. Taking expectations and applying the standard quadratic-form moment identity ($\mathbb{E}[X^\top A X] = \operatorname{tr}(A\Sigma) + \mu^\top A\, \mu$ for $X \sim (\mu, \Sigma)$ with $\Sigma = \operatorname{Cov}(X, X^\top)$) gives

$$\mathbb{E}[\sigma_T \cdot T] \;=\; \operatorname{tr}\big(\mathcal{M}\,\Sigma\big) \;+\; \mu^\top \mathcal{M}\, \mu, \tag{26}$$

which is the master identity of the appendix. Two structural simplifications recur:

$$\mathcal{M}\,\mathbf{1} \;=\; 0, \qquad \operatorname{tr}(\mathcal{M}\,\Sigma) \;=\; \operatorname{tr}(\Sigma) - (n+1)^{-1} \mathbf{1}^\top \Sigma\, \mathbf{1}. \tag{27}$$

The first kills the bias term $\mu^\top \mathcal{M}\, \mu$ whenever μ_j is constant in j; the second reduces the trace to a weighted-sum expression that is family-tractable.

A.2 GBM family

Under (9) at $\mu = 0$, the marginals are $X_j = x_0 \exp(L_j)$ with $L_j \sim \mathcal{N}(-\sigma^2 t_j/2, \sigma^2 t_j)$; the martingale pin $\mathbb{E}[X_j] = x_0$ keeps μ_j constant in j, so the bias term in (26) vanishes. The covariance kernel of (11) gives

$$\Sigma_{jk}^{\text{GBM}} \;=\; x_0^2\big(\exp(\sigma^2 \min(t_j, t_k)) - 1\big), \tag{28}$$

and the trace identity reduces to

$$\mathbb{E}[\sigma_T \cdot T]^{\text{GBM}} \;=\; \operatorname{tr}\big(\mathcal{M}\,\Sigma^{\text{GBM}}\big) \;=\; \sum_{j=0}^{n} \Sigma_{jj}^{\text{GBM}} \;-\; (n+1)^{-1} \sum_{j,k=0}^{n} \Sigma_{jk}^{\text{GBM}}. \tag{29}$$

Evaluating at the canonical pin ($x_0 = 4000, \sigma = 0.10/\sqrt{12}, T = 12, n = 5000$) yields $\mathbb{E}[\sigma_T]^{\text{GBM}} = 1.117130{\times}10^7$, which matches the hand-derivation script `scratch/2026-05-13-task-4.2-discrete-moments/derivation.py` to 13 significant figures.

A.3 OU family (stationary at $x_0 = \bar{\mu}$)

Under (14) the mean dynamics $\mathbb{E}[X_j] = \bar{\mu} + (x_0 - \bar{\mu})e^{-\theta t_j}$ collapse to $\mathbb{E}[X_j] = \bar{\mu}$ at the canonical pin $x_0 = \bar{\mu}$; again μ_j is constant in j and the bias term vanishes. The Vasicek covariance kernel of (16) gives

$$\Sigma_{jk}^{\text{OU}} \;=\; \frac{\sigma^2}{2\theta}\, e^{-\theta|t_k - t_j|}\big(1 - e^{-2\theta \min(t_j,t_k)}\big). \tag{30}$$

For $\theta T \gg 1$ (canonical horizon ≈ 17 relaxation times) the transient factor $(1 - e^{-2\theta \min(t_j,t_k)}) \to 1$ on the bulk of the grid, and the kernel approaches its stationary $\sigma^2/(2\theta) \cdot e^{-\theta|t_k - t_j|}$ form. The trace identity reduces to

$$\mathbb{E}[\sigma_T \cdot T]^{\text{OU}} \;=\; \operatorname{tr}\big(\mathcal{M}\,\Sigma^{\text{OU}}\big) \;\xrightarrow[\theta T \to \infty]{}\; (n+1) \cdot \frac{\sigma^2}{2\theta} \;-\; (n+1)^{-1} \sum_{j,k} \frac{\sigma^2}{2\theta}\, e^{-\theta|t_k - t_j|}, \tag{31}$$

whose continuous-time limit per T is exactly $\sigma^2/(2\theta)$ — the one-line closed form of (18). This is the unique family in this paper where the T-dependence collapses to a constant in the stationary regime.

A.4 Merton family (non-constant μ_j)

Under (19) the mean dynamics carry a drift correction from the compound-Poisson jump:

$$\mathbb{E}[X_j]^{\text{MERTON}} \;=\; x_0\, \exp\!\big(\lambda(\mathbb{E}[e^J] - 1)\, t_j\big), \tag{32}$$

where $J \sim \mathcal{N}(\mu_J, \sigma_J^2)$ is the log-jump amplitude and $\mathbb{E}[e^J] = \exp(\mu_J + \sigma_J^2/2)$. Unlike GBM and OU, μ_j is NOT constant in j under the canonical martingale pin (which sets the drift compensator $\alpha = -\lambda(\mathbb{E}[e^J] - 1)$ on the diffusion side, NOT on the mean): the bias term $\mu^\top \mathcal{M}\,\mu$ in (26) no longer vanishes and must be carried explicitly. The covariance kernel decomposes into diffusion and jump contributions:

$$\Sigma_{jk}^{\text{MERTON}} \;=\; \mathbb{E}[X_j]\mathbb{E}[X_k]\Big(\exp\!\big((\sigma^2 + \lambda\,\mathbb{E}[e^{2J}] - \lambda\,(\mathbb{E}[e^J])^2)\min(t_j, t_k)\big) - 1\Big). \tag{33}$$

The discrete trace identity (26) now reads

$$\mathbb{E}[\sigma_T \cdot T]^{\text{MERTON}} \;=\; \operatorname{tr}\big(\mathcal{M}\,\Sigma^{\text{MERTON}}\big) \;+\; \sum_{j=0}^{n} \mu_j^2 \;-\; (n+1)^{-1}\Big(\sum_{j=0}^{n} \mu_j\Big)^2, \tag{34}$$

where the trailing two terms expand the non-vanishing $\mu^\top \mathcal{M}\, \mu$ bias contribution. The closed-form evaluation at the canonical pin ($x_0 = 4000, \sigma = 0.10/\sqrt{12}, \lambda = 0.5, \mu_J = -0.05, \sigma_J = 0.10, T = 12, n = 5000$) yields $\mathbb{E}[\sigma_T]^{\text{MERTON}}$ matching the hand-derivation script to 12 significant figures. The Wave-2 MQ independent re-derivation (§7.5) reports three-way agreement (MQ-independent, paper closed-form, code analytic) at relative error $\leq 5.95 \times 10^{-12}$.

B Anti-fishing audit trail across spec amendments

The plan docs/plans/2026-05-11-stochastic-fx-variant.md carried four spec amendments through Wave-1 dispute and Wave-2 plan-exit review. Each amendment surfaced a concrete fishing risk at the implementation phase; each was routed through CORRECTIONS-α with the three concrete anti-fishing prohibitions preserved invariant. The prohibitions, in canonical phrasing:

- **(NF-1) No silent threshold relaxation.** The three Phase A/B/C tolerances (NUMERICAL_IDENTITY_TOL 10^{-6}, MOMENT_REL_TOL $= 0.05$, KS_PVALUE_FLOOR $= 0.01$) are PINNED and invariant across all four families.
- **(NF-2) No seed re-roll.** The canonical seed $= 42$ is PINNED; on canonical-pin failure, route to CORRECTIONS-α (re-derive analytic or refine discretization), NEVER to a different seed.
- **(NF-3) No N-floor lowering.** $N = 1000$ is the canonical N-floor and the source of the Type-I budget; tightening Phase B tolerances post-hoc to engineer PASS at $N = 1000$ is silent-fishing.

The four amendments, in plan-order:

- **§16.3 — Task 3.1 grid-density correction.** Implementation discovered a grid-density / tolerance mismatch in the trivial-degenerate test class: the coarse-grid + tight-threshold pairing is canonical for these tests, against an initial implementation that paired coarse-grid + loose-threshold. Resolution: pin coarse-grid + tight-threshold as canonical, document the rationale. (NF-1) preserved (no threshold change).
- **§16.4 — Task 3.3 σ_T statistic mismatch.** Implementation found that the moment-matching test originally scoped to Task 3.3 was a duplicate of work performed under Task 4.2 Phase B. Resolution: route the moment-matching to Phase B exclusively, eliminating the duplication. (NF-1) and (NF-3) preserved (no relaxation, no N change).
- **§16.5 — Task 4.2 mean-only Phase B amendment.** Implementation discovered that the Isserlis-Gaussian variance formula carries an MC-noise floor of 8%–30% at $N = 1000$, structurally incompatible with the 5% MOMENT_REL_TOL. Resolution: amend Phase B to mean-only (spec v0.5 §11.6); record variance rel-err as audit-trail only; disclose the family-dependent Type-I rates (0%/7%/30% at OU/GBM/Merton). (NF-1), (NF-2), (NF-3) all preserved: the threshold is unchanged, the seed is unchanged, N is unchanged. The Type-I budget is the cost of the N-floor invariant; v0.5 §11.6 discloses it honestly.
- **§16.7 — Task 4.2 per-family Phase C dispatch.** Implementation discovered that the moment-matched lognormal Phase C reference, generalized from GBM to all families, returns $p = 3.41 \times 10^{-21}$ at Merton — 19 orders of magnitude below the 0.01 floor. Resolution: per-family dispatch on the reference *shape* (lognormal for GBM, gamma for OU, empirical-CDF from a pinned $N_{\text{REF}} = 100{,}000$ auxiliary run for Merton, spec v0.5 §11.7). (NF-1) preserved

(KS_PVALUE_FLOOR = 0.01 invariant across all three shapes); the dispatch is on geometric content of the marginals, not on threshold tuning. The five anti-fishing sub-prohibitions of the disposition (single threshold, N_{REF} pinned, $N_{\text{REF-SEED}}$ pinned, reference is independently sampled not `.fit()`'d, PIN Z1.4 interpretation unchanged) are itemized in Remark 5.3.

The Wave-1 dispute over each amendment was independently re-reviewed by RC and MQ in scratch/2026-05-11-stochastic-fx-spec-review/wave-1-v0.4/ and scratch/2026-05-11-stochastic-fx-spec-review/wave-1-v0.5/; the Wave-2 plan-exit verdict at scratch/2026-05-11-stochastic-fx-spec-review/wave-2/rc-verdict.md and scratch/2026-05-11-stochastic-fx-spec-review/wave-2/mq-verdict.md records anti-fishing prohibitions PRESERVED across all four dispositions. The joint plan-exit verdict is ACCEPT_WITH_FLAGS (two informational flags carried forward, none anti-fishing-relevant).

C cadCAD-implementation spec (forward-looking)

This appendix sketches the cadCAD implementation spec that this paper nominates for the immediate next deliverable. The sketch is forward-looking; no implementation work is performed in this paper.

Directory structure. A new sub-package simulations/cadcad/ under the existing simulations/ functional-python three-tier layout (types/ value, modules/ callable, utils/ I/O boundary), with the same tier-import discipline that governs simulations/stochastic_fx/:

- simulations/cadcad/types.py — frozen-dataclass containers for the six state-variable values, the simulation_config_dict, and the `SimulationVerdict` value type.
- simulations/cadcad/state.py — initial-state constructors, PSUB-sequence composition.
- simulations/cadcad/psubs.py — six PSUB implementations (one per role in §6.2); callable-tier frozen dataclasses with `__call__`.
- simulations/cadcad/runner.py — boundary layer: invokes cadCAD [7] engine, materialises run logs to Parquet, computes audit-block hashes.
- simulations/cadcad/tests/ — pytest suite crossing the four FX families and the cost-factor trajectory.

State-variable declarations (cadCAD config). A schematic of the simulation_config_dict entries (implementing the signature of §6.2):

```
initial_state = {
    "X": x_0,           # FX rate, R_+
    "S": 0.0,           # sigma_T accumulator, R_+
    "C": C_0,           # cost trajectory, R_+
    "P": 0.0,           # replicating-portfolio MtM, R
    "W": W_0,           # wage-funded premium balance, R_+
    "R": 0.0,           # ratchet accumulator, R
}
sys_params = {
    "family": ["gbm", "ou", "merton", "deterministic"],
    "seed":   [42],
    "N":      [1000],
```

```
    # ... per-family parameter pin
}
partial_state_update_blocks = [
    psub_fx_advance,        # PSUB-1
    psub_sigma_t_accum,     # PSUB-2
    psub_cost_advance,      # PSUB-3
    psub_portfolio_update,  # PSUB-4
    psub_premium_drain,     # PSUB-5
    psub_ratchet_accum,     # PSUB-6
]
```

PSUB implementations. Each PSUB is a frozen-dataclass callable per the functional-python skill: stateless transforms with explicit input and output state-variable signatures. PSUB-1 dispatches to the verified generators of simulations/stochastic_fx/generators.py; PSUB-2 to simulations/stochastic_fx/variance_proxy.py; PSUB-4 to the frozen 12-leg IronCondor strip at simulations/saas_builder/data/IronCondor_strip.json (PIN Z1.6). PSUB-3, PSUB-5, PSUB-6 are new and live entirely in simulations/cadcad/psubs.py.

Integration tests. The test suite crosses the four FX families (GBM / OU / Merton / deterministic) with at least two cost-trajectory parameter pins (low-volatility AI-cost demand, high-volatility AI-cost demand) for a $4 \times 2 = 8$ test matrix. Per-cell invariants:

- (INV-1) state-variable types and value-range non-negativity where applicable;
- (INV-2) audit-block hash bit-equality on intra-process re-run, inheriting PIN Z1.2;
- (INV-3) PSUB-2 σ_T accumulator output agrees with the verified stochastic_fx output bit-for-bit;
- (INV-4) PIN Z1.6 byte-equality re-check on the IronCondor strip JSON;
- (INV-5) cross-process determinism re-check following the Wave-2 RC protocol of §7.4.

Anti-fishing invariants carried forward. The three concrete anti-fishing prohibitions of Appendix B carry forward unchanged: tolerances (NF-1), seed (NF-2), and N-floor (NF-3) are PINNED for the cadCAD wrapping at the same canonical values as the stochastic-FX verification floor. The wrapping introduces no new threshold, seed, or N knob; any future amendment routes through CORRECTIONS-α per the standard plan-amendment protocol.

Doc-first sequencing. The cadCAD-spec (docs/specs/2026-05-16-cadcad-wrapping-design.md, to be written) is the next deliverable, gated on this paper's compile-clean commit. Implementation follows the spec; no simulations/cadcad/ code is added before the spec lands.

References

[1] Hull, John C. (2017). *Options, Futures, and Other Derivatives*, 10th ed. Pearson. §15.

[2] Karatzas, Ioannis & Steven E. Shreve (1991). *Brownian Motion and Stochastic Calculus*, 2nd ed. Springer. §5.6 (Ornstein–Uhlenbeck process).

[3] Andersen, Leif B. G. & Vladimir V. Piterbarg (2010). *Interest Rate Modeling, Volume I: Foundations and Vanilla Models*. Atlantic Financial Press. §2.7 (jump-diffusion moments).

[4] Carr, Peter & Dilip Madan (2001). Optimal Positioning in Derivative Securities. *Quantitative Finance* 1(1), 19–37.

[5] Merton, Robert C. (1976). Option Pricing When Underlying Stock Returns Are Discontinuous. *Journal of Financial Economics* 3, 125–144.

[6] Vasicek, Oldrich (1977). An Equilibrium Characterization of the Term Structure. *Journal of Financial Economics* 5(2), 177–188.

[7] cadCAD: complex Adaptive Dynamics Computer-Aided Design. https://github.com/cadCAD-org/cadCAD.

[8] Panoptic: Perpetual Options on Uniswap. Technical whitepaper. 2023.

[9] Abrigo Analytics (2026). Stochastic-FX variant design spec, v0.5. `docs/specs/2026-05-11-stochastic-fx-variant-design.md`, in `abrigo-analytics` repository.

[10] Abrigo Analytics (2026). Stochastic-FX variant implementation plan, v0.6. `docs/plans/2026-05-11-stochastic-fx-variant.md`, in `abrigo-analytics` repository.

[11] Abrigo Analytics (2026). AI-Cost Factor Model design spec, v0.2.1 (under construction). `docs/specs/2026-05-16-ai-cost-factor-model-design.md`, in `abrigo-analytics` repository.